Current agent behavior

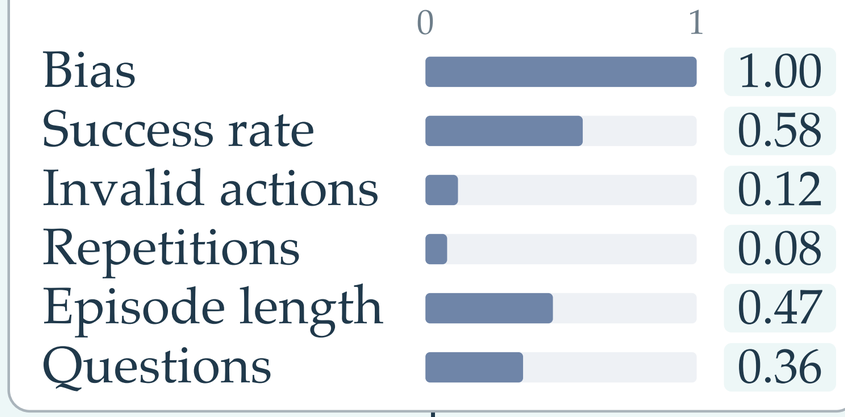

x_k

BATS: select one regime

$$m_k = \arg\max_m \left[\widehat{\mu}_{m,k}(x_k) + \alpha\, \sigma_{m,k}(x_k)\right]$$

DAgger | **PPO**

Update the shared agent

$\theta_k \longrightarrow \theta_{k+1}$

Evaluate updated agent

x_{k+1}, $\mathrm{SR}_{\mathrm{after}}$

Update bandit estimates

$\rho_k = \mathrm{SR}_{\mathrm{after}} - \mathrm{SR}_{\mathrm{before}}$
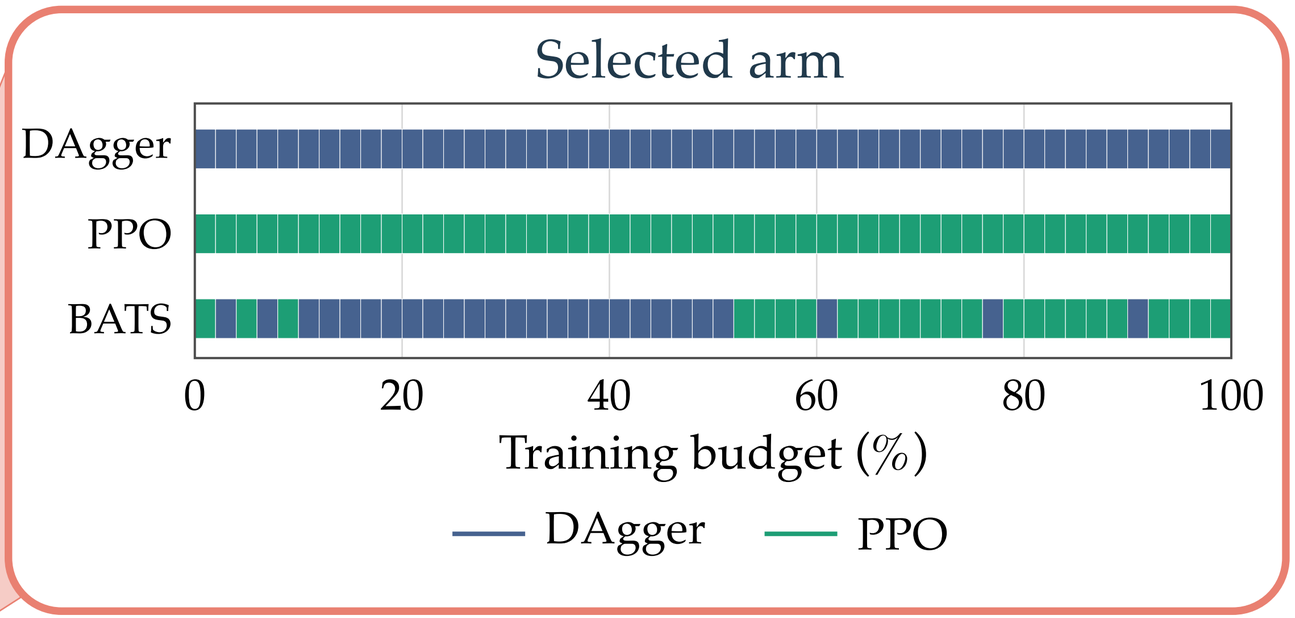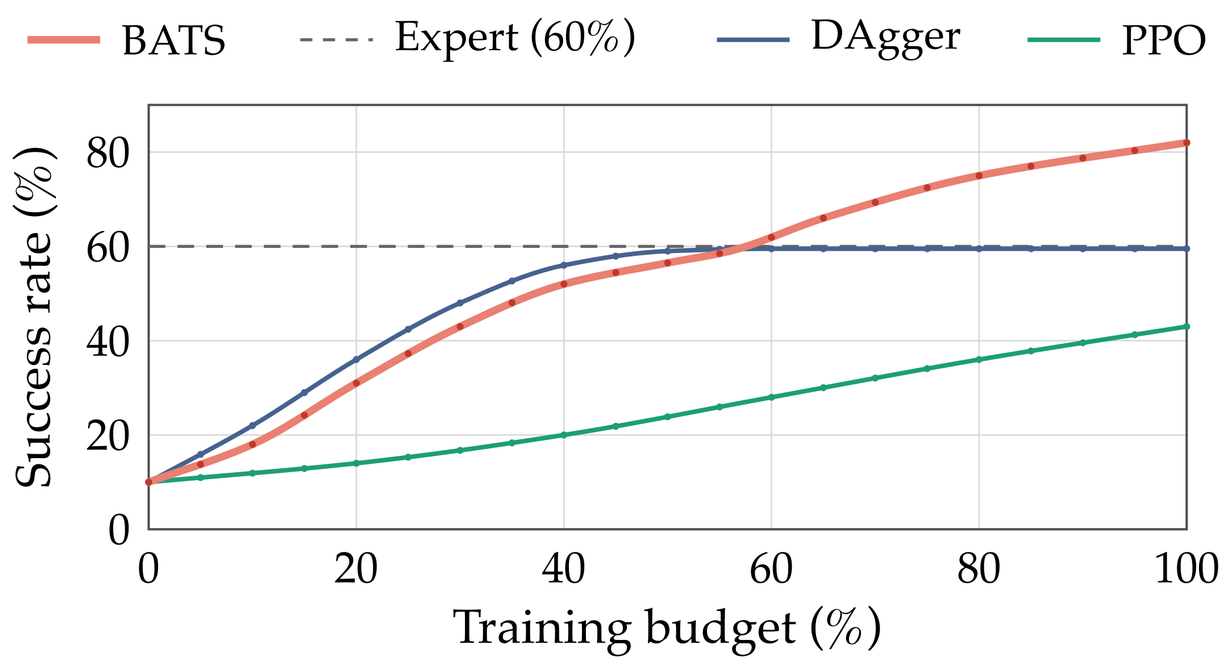
$(\widehat{\mu}, \sigma)_k \longrightarrow (\widehat{\mu}, \sigma)_{k+1}$